Question 21 (Single Topic)

You have deployed multiple versions of an image classification model on AI Platform. You want to monitor the performance of the model versions over time. How should you perform this comparison?

A. Compare the loss performance for each model on a held-out dataset.
B. Compare the loss performance for each model on the validation data.
C. Compare the receiver operating characteristic (ROC) curve for each model using the What-If Tool.
D. Compare the mean average precision across the models using the Continuous Evaluation feature.

Expose Correct Answer

Question 22 (Single Topic)

You trained a text classification model. You have the following SignatureDefs:

You started a TensorFlow-serving component server and tried to send an HTTP request to get a prediction using: headers = {"content-type": "application/json"} json_response = requests.post('http: //localhost:8501/v1/models/text_model:predict', data=data, headers=headers)
What is the correct way to write the predict request?

A. data = json.dumps({â€signature_nameâ :€â€seving_defaultâ ,€â€instancesâ€ [[â€˜abâ ,™€â€˜bcâ ,™€â€˜cdâ€™]]})
B. data = json.dumps({â€signature_nameâ :€â€serving_defaultâ ,€â€instancesâ€ [[â€˜aâ ,™€â€˜bâ ,™€â€˜câ ,™€â€˜dâ ,™€â€˜eâ ,™€â€˜fâ€™]]})
C. data = json.dumps({â€signature_nameâ :€â€serving_defaultâ ,€â€instancesâ€ [[â€˜aâ ,™€â€˜bâ ,™€â€˜câ€™], [â€˜dâ ,™€â€˜eâ ,™€â€˜fâ€™]]})
D. data = json.dumps({â€signature_nameâ :€â€serving_defaultâ ,€â€instancesâ€ [[â€˜aâ ,™€â€˜bâ€™], [â€˜câ ,™€â€˜dâ€™], [â€˜eâ ,™€â€˜fâ€™]]})

Expose Correct Answer

Question 23 (Single Topic)

Your organization's call center has asked you to develop a model that analyzes customer sentiments in each call. The call center receives over one million calls daily, and data is stored in Cloud Storage. The data collected must not leave the region in which the call originated, and no Personally Identifiable Information (PII) can be stored or analyzed. The data science team has a third-party tool for visualization and access which requires a SQL ANSI-2011 compliant interface. You need to select components for data processing and for analytics. How should the data pipeline be designed?

A. 1= Dataflow, 2= BigQuery
B. 1 = Pub/Sub, 2= Datastore
C. 1 = Dataflow, 2 = Cloud SQL
D. 1 = Cloud Function, 2= Cloud SQL

Expose Correct Answer

Question 24 (Single Topic)

You are an ML engineer at a global shoe store. You manage the ML models for the company's website. You are asked to build a model that will recommend new products to the user based on their purchase behavior and similarity with other users. What should you do?

A. Build a classification model
B. Build a knowledge-based filtering model
C. Build a collaborative-based filtering model

D. Build a regression model using the features as predictors

Expose Correct Answer

Question 25 (Single Topic)

You work for a social media company. You need to detect whether posted images contain cars. Each training example is a member of exactly one class. You have trained an object detection neural network and deployed the model version to AI Platform Prediction for evaluation. Before deployment, you created an evaluation job and attached it to the AI Platform Prediction model version. You notice that the precision is lower than your business requirements allow. How should you adjust the model's final layer softmax threshold to increase precision?

A. Increase the recall.
B. Decrease the recall.
C. Increase the number of false positives.
D. Decrease the number of false negatives.

Expose Correct Answer