Custom View Settings

Question #21 *Topic 1*

You have deployed multiple versions of an image classification model on AI Platform. You want to monitor the performance of the model versions over time. How should you perform this comparison?

A. Compare the loss performance for each model on a held-out dataset.

B. Compare the loss performance for each model on the validation data.

C. Compare the receiver operating characteristic (ROC) curve for each model using the What-If Tool.

D. Compare the mean average precision across the models using the Continuous Evaluation feature.

Question #22 *Topic 1*

You trained a text classification model. You have the following SignatureDefs:

```
signature_def['serving_default']:
  The given SavedModel SignatureDef contains the following input(s):
    inputs['text'] tensor_info:
        dtype: DT_STRING
        shape: (-1, 2)
        name: serving_default_text: 0
  The given SavedModel SignatureDef contains the following output(s):
    outputs ['Softmax'] tensor_info:
        dtype: DT_FLOAT
        shape: (-1, 2)
        name: StatefulPartitionedCall:0
  Method name is: tensorflow/serving/predict
```

You started a TensorFlow-serving component server and tried to send an HTTP request to get a prediction using: headers = {"content-type": "application/json"} json_response = requests.post('http: //localhost:8501/v1/models/text_model:predict', data=data, headers=headers)
What is the correct way to write the predict request?

A. data = json.dumps({ג€signature_nameג :€ג€seving_defaultג ,€ג€instancesג€ [['ab', 'bc', 'cd']]})

B. data = json.dumps({ג€signature_nameג :€ג€serving_defaultג ,€ג€instancesג€ [['a', 'b', 'c', 'd', 'e', 'f']]})

C. data = json.dumps({ג€signature_nameג :€ג€serving_defaultג ,€ג€instancesג€ [['a', 'b', 'c'], ['d', 'e', 'f']]})

D. data = json.dumps({ג€signature_nameג :€ג€serving_defaultג ,€ג€instancesג€ [['a', 'b'], ['c', 'd'], ['e', 'f']]})

Question #23 *Topic 1*

Your organization's call center has asked you to develop a model that analyzes customer sentiments in each call. The call center receives over one million calls daily, and data is stored in Cloud Storage. The data collected must not leave the region in which the call originated, and no Personally Identifiable Information (PII) can be stored or analyzed. The data science team has a third-party tool for visualization and access which requires a SQL ANSI-2011 compliant interface. You need to select components for data processing and for analytics. How should the data pipeline be designed?

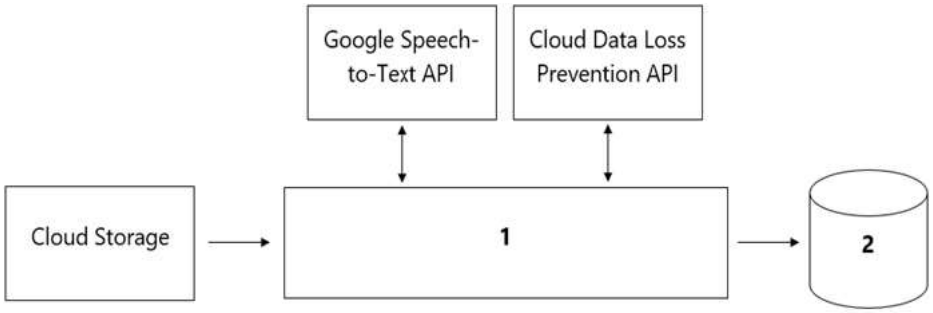

A. 1= Dataflow, 2= BigQuery

B. 1 = Pub/Sub, 2= Datastore

C. 1 = Dataflow, 2 = Cloud SQL

D. 1 = Cloud Function, 2= Cloud SQL

Question #24 *Topic 1*

You are an ML engineer at a global shoe store. You manage the ML models for the company's website. You are asked to build a model that will recommend new products to the user based on their purchase behavior and similarity with other users. What should you do?

A. Build a classification model

B. Build a knowledge-based filtering model

C. Build a collaborative-based filtering model

D. Build a regression model using the features as predictors

Previous Questions

Next Questions